

**STAT 515**  
**Homework #8 WITH SOLUTIONS**

1. The matrix exponential function is defined as

$$\exp\{M\} = \sum_{i=0}^{\infty} \frac{M^i}{i!}.$$

However, this definition does not provide a suitable method for calculating  $\exp\{M\}$  for a given  $M$ . One simple alternative is to use one of the two formulas

$$\exp\{M\} = \lim_{n \rightarrow \infty} \left(I + \frac{M}{n}\right)^n = \lim_{n \rightarrow \infty} \left[\left(I - \frac{M}{n}\right)^{-1}\right]^n.$$

For the post office problem on HW #7, the rate matrix is given by

$$R = \begin{bmatrix} -8 & 8 & 0 & 0 & 0 \\ 12 & -20 & 8 & 0 & 0 \\ 0 & 12 & -20 & 8 & 0 \\ 0 & 0 & 12 & -20 & 8 \\ 0 & 0 & 0 & 12 & -12 \end{bmatrix},$$

where rates are in hours. Approximate the value of  $\exp\{.5R\}$ , which is the transition probability matrix for a time step of 30 minutes, using two methods:

- (a) For successively larger powers of 2, i.e.,  $n = 2, 4, 8, \dots$ , find the value of  $(I + .5R/n)^n$ . Continue until the change in each entry is smaller than  $10^{-5}$ . Report your final value of  $n$  and your final approximation of  $\exp\{.5R\}$ .

**Solution:** First, let's set up the  $R$  matrix:

```
> R <- matrix(c(-8, 12, 0, 0, 0, 8, -20, 12, 0, 0, 0, 8, -20,
+              12, 0, 0, 0, 8, -20, 12, 0, 0, 0, 8, -12), 5, 5)
```

Next, we'll use a loop to continue trying larger and larger  $n$ , starting with  $n = 2$  and doubling it at every step, until the maximum change is smaller than  $10^{-5}$ :

```
> lastAnswer <- matrix(0, 5, 5)
> finished <- FALSE
> k <- 0
> n <- 1
> while (!finished) {
+   k <- k+1
+   n <- n*2 # double the n
+   answer <- diag(5) + .5*R/n
+   for (j in 1:k) { # raise answer to the nth power by squaring repeatedly
+     answer <- answer %*% answer
+   }
+   finished <- all(abs(answer - lastAnswer) < 1e-5) # check for convergence
+   lastAnswer <- answer
+ }
> lastAnswer

      [,1]      [,2]      [,3]      [,4]      [,5]
[1,] 0.4176125 0.2665951 0.1623088 0.09580071 0.05768283
[2,] 0.3998927 0.2611831 0.1668329 0.10513201 0.06695930
[3,] 0.3651949 0.2502494 0.1754179 0.12357079 0.08556707
[4,] 0.3233274 0.2365470 0.1853562 0.14607047 0.10869892
[5,] 0.2920193 0.2259876 0.1925259 0.16304838 0.12641878
> n
[1] 32768
```

So the desired accuracy required  $n = 2^{15} = 32,768$ .

- (b) Repeat the same procedure as in part (a) but use  $[(I - .5R/n)^{-1}]^n$  instead.

**Solution:** This requires only minor modifications to the code above:

```
> lastAnswer <- matrix(0, 5, 5)
> finished <- FALSE
> k <- 0
> n <- 1
> while (!finished) {
+   k <- k+1
+   n <- n*2 # double the n
+   answer <- solve(diag(5) - .5*R/n)
+   for (j in 1:k) { # raise answer to the nth power by squaring repeatedly
+     answer <- answer %**% answer
+   }
+   finished <- all(abs(answer - lastAnswer) < 1e-5) # check for convergence
+   lastAnswer <- answer
+ }
> lastAnswer
```

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	0.4176172	0.2665964	0.1623075	0.09579832	0.05768059
[2,]	0.3998946	0.2611839	0.1668326	0.10513091	0.06695803
[3,]	0.3651919	0.2502489	0.1754190	0.12357216	0.08556805
[4,]	0.3233193	0.2365445	0.1853582	0.14607473	0.10870316
[5,]	0.2920080	0.2259833	0.1925281	0.16305474	0.12642582

```
> n
[1] 32768
```

This example also required  $n = 2^{15} = 32,768$ .

- (c) Use the `expm` function in R or Matlab to evaluate  $\exp\{.5R\}$  and compare with the two approximations you obtained.

*In R, you will have to install and load the package called `Matrix`. Do this using `install.packages("Matrix")` and then `library(Matrix)`.*

**Solution:**

```
> library(Matrix)
> expm(.5*R)
5 x 5 Matrix of class "dgeMatrix"
      [,1] [,2] [,3] [,4] [,5]
[1,] 0.4176149 0.2665957 0.1623082 0.09579952 0.05768171
[2,] 0.3998936 0.2611835 0.1668328 0.10513146 0.06695866
[3,] 0.3651934 0.2502491 0.1754184 0.12357147 0.08556756
[4,] 0.3233234 0.2365458 0.1853572 0.14607260 0.10870104
[5,] 0.2920136 0.2259855 0.1925270 0.16305156 0.12642230
```

The answers in parts (a) and (b) are very close to the `expm` answer.

2. Problem 3 in homework #7 described two video game machines at an amusement park. For video game  $i$ , each period when it is being used is exponentially distributed with mean  $1/\alpha_i$  hours and each period when it is not being used is exponentially distributed with mean  $1/\beta_i$  hours, independent of the other machine. Furthermore,  $\alpha_1 = 2$ ,  $\alpha_2 = 3$ ,  $\beta_1 = 5$ , and  $\beta_2 = 6$ .

- (a) If neither machine is in use when the park opens at 8:00am, find the probability that both machines are in use at 9:30am.

**Solution:** To find the answer, we can find the probability transition matrix  $P(1.5)$ , which is given by the matrix exponential  $\exp\{1.5R\}$ :

```
> R <- matrix(c(-11, 3, 2, 0, 6, -8, 0, 2, 5, 0, -8, 3, 0, 5, 6, -5), 4, 4)
> expm(1.5*R)
```

```

4 x 4 Matrix of class "dgeMatrix"
      [,1] [,2] [,3] [,4]
[1,] 0.09524491 0.1904890 0.2380893 0.4761767
[2,] 0.09524452 0.1904894 0.2380884 0.4761777
[3,] 0.09523573 0.1904707 0.2380985 0.4761951
[4,] 0.09523534 0.1904711 0.2380975 0.4761960

```

The answer is the probability that starting in the first state at time zero, the chain is in the fourth state at time 1.5:

```

> expm(1.5*R)[1,4]
[1] 0.4761767

```

- (b) Simulate 10,000 realizations of the Markov chain and give a 95% confidence interval for the probability in part (a) based on your simulation. Does your empirical estimate agree with the theoretical value?

**Solution:** I will make some minor changes to the code used in homework #7, problem 3(c):

```

> n <- 10000
> finalState <- rep(0, n)
> maxTime <- 1.5 # This is the cutoff time.
> for (count in 1:n) {
+   currentTime <- 0
+   currentState <- 1 # Assume that we always start in state 1 at time 0
+   finished <- FALSE # We'll set this to TRUE when it's time to stop.
+   while (!finished) {
+     deltaTime <- rexp(1, rate = -R[currentState, currentState])
+     if (currentTime + deltaTime > maxTime) {
+       # Now we need to finish this chain and declare the final state decided
+       finalState[count] <- currentState
+       deltaTime <- maxTime - currentTime
+       finished <- TRUE
+     }
+     currentTime <- currentTime + deltaTime
+     possibleMoves <- (1:4)[-currentState]
+     currentState <- sample(possibleMoves, 1, prob=R[currentState,possibleMoves])
+   }
+ }
> table(finalState) / n
finalState
      1      2      3      4
0.0986 0.1886 0.2488 0.4640

```

Notice how close the four probabilities are to the first row (really, all rows) of the matrix found in part (a). Here is a 95% confidence interval for the probability of ending in the fourth state:

```

> phat <- sum(finalState==4) / n
> phat + c(-1.96, 1.96) * sqrt(phat * (1-phat) / n)
[1] 0.4542254 0.4737746

```

3. Suppose that  $X_1, X_2, X_3, X_4$  are i.i.d. from a uniform(0,1) distribution. Let  $S = X_1 + X_2 + X_3 + X_4$ .

- (a) Find  $P(S < 1)$  exactly using a four-dimensional integral. (Hint: This is not too difficult.)

**Solution:** The joint density on the four-dimensional hypercube  $(0,1)^4$  is just the constant 1.

The probability of  $S < 1$  is found by integrating this density over the region where  $S < 1$ :

$$\begin{aligned}
 \int_0^1 \int_0^{1-w} \int_0^{1-w-x} \int_0^{1-w-x-y} dz \, dy \, dx \, dw &= \int_0^1 \int_0^{1-w} \int_0^{1-w-x} (1-w-x-y) \, dy \, dx \, dw \\
 &= \frac{1}{2} \int_0^1 \int_0^{1-w} \frac{1}{2} (1-w-x)^2 \, dx \, dw \\
 &= \frac{1}{6} \int_0^1 (1-w)^3 \, dw \\
 &= \frac{1}{24}.
 \end{aligned}$$

- (b) Now consider  $P(S < 1.5)$ . This is much more difficult to find analytically. Instead, use Monte Carlo simulation to approximate this probability. Give a 99% confidence interval for the true probability, and use a large enough sample so that your interval is no wider than 0.01. Report the sample size you used in addition to the interval.

**Solution:** Since we are finding a proportion, and the standard deviation of a sample proportion is not more than  $1/\sqrt{4n}$ , we could take  $n$  large enough so that  $2.58/\sqrt{4n} < .005$  since the 99% confidence interval is  $\hat{p} \pm 2.58 \times (\text{standard error})$ . This leads to  $n$  larger than about 66,000. Let's use 100,000 just for a nice round number:

```

> n <- 1e6
> x <- matrix(runif(4*n), ncol=4)
> phat <- sum(rowSums(x)<1.5) / n
> phat + c(-2.58, 2.58) * sqrt(phat * (1-phat) / n)
[1] 0.1995289 0.2015951

```

So the 99% confidence interval obtained from a sample of size 100,000 has a width of much less than 0.01 (the width is about 0.002).

- (c) The central limit theorem approximation to  $P(S < 1.5)$  is  $P(Y < 1.5)$ , where  $Y$  is a normal random variable with the same mean and variance as  $S$ . Based on your answer to part (b), how good does the central limit theorem approximation appear in this case?

**Solution:** Since  $X_i$  has mean  $1/2$  and variance  $1/12$ ,  $S$  has mean 2 and variance  $1/3$ . Thus, the central limit theorem approximation is

```

> pnorm(1.5, mean=2, sd=1/sqrt(3))
[1] 0.1932381

```