

Typical intro. to prob. theory + stat. inference: $\begin{pmatrix} \text{Stat} \\ S13+ \\ S14 \end{pmatrix}$
 most problems are solved analytically e.g. using tools from multivariate calculus.

In [Above topics: required background for this class]
 In practice: only a small number of problems admit analytical solutions.

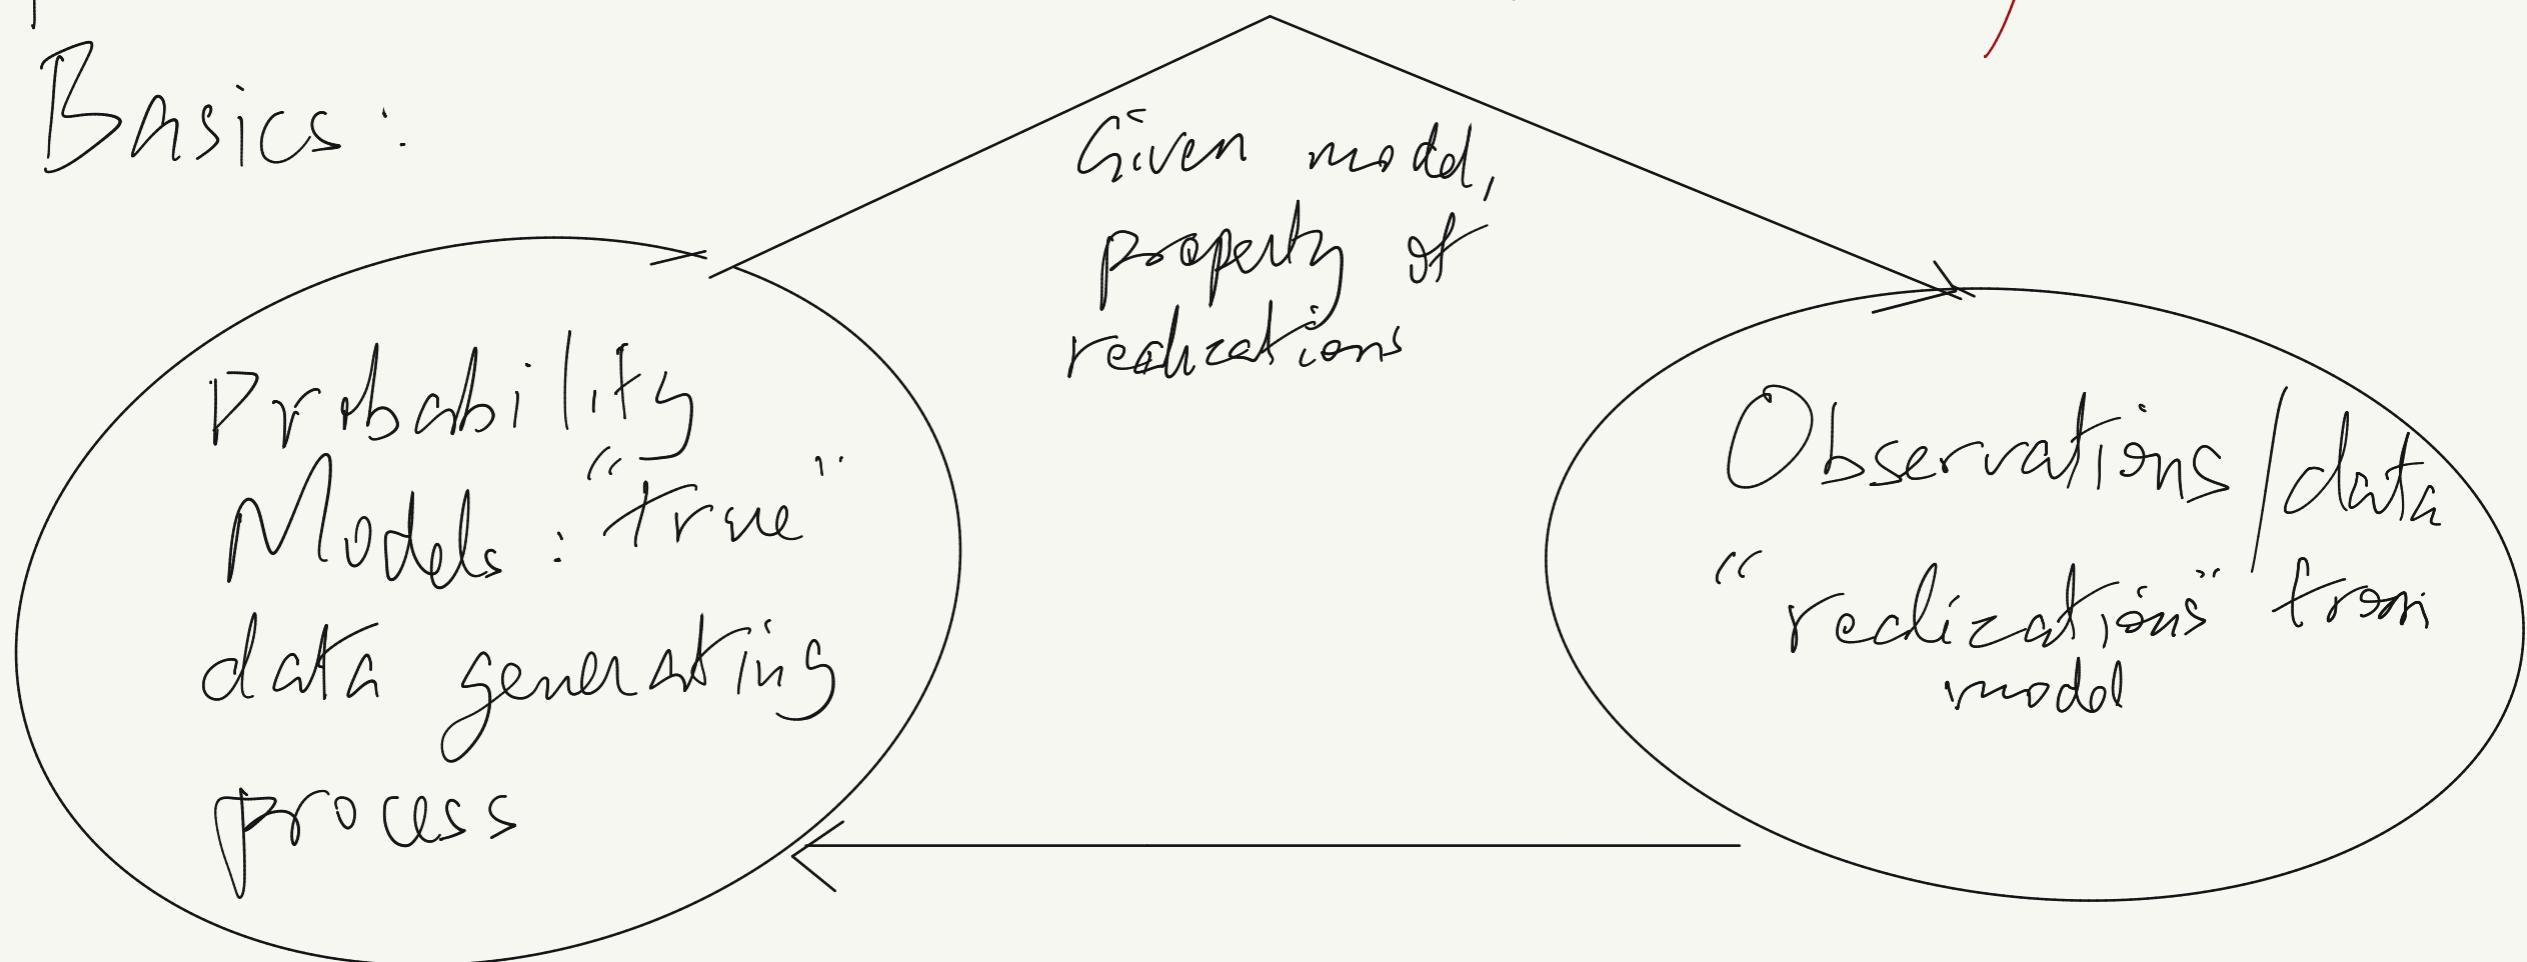
Most require computer algorithms.

This is the subject of this course: learn computer algorithms in order to solve probability and stat. inf. problems.

Probability

computing
mostly
simulation/Monte Carlo

Basics:



Statistical Inference

learn about model based

on data Computing: optimization; also simulation/Monte Carlo
 e.g. Bayesian inf., hyp. testing,
 goodness of fit

Of course, prob. is central to stat. inf.: without understanding model, cannot carry out inference. e.g. need joint distr. to obtain likelihood fn. e.g. properties of fitted prob. model: used to assess goodness of fit of model to data.

Probability calculations can be analytically intractable but easy w/ computer algorithms.

Simple e.g. $\underline{X} \sim N_d(\underline{\mu}, \Sigma)$. X_{\max}, X_{\min} : max and min of vector \underline{X} . For some $\underline{\mu}, \Sigma$, what is $M = P(X_{\max}/X_{\min} > c)$?

Analytical form of X_{\max}/X_{\min} complicated.

E.g. X_{\max}, X_{\min} distr. for bivariate Normal
(Nadarajah & Kotz, 2008)

But Monte Carlo approximation is easy:

Simulate $\underline{X}^{(1)}, \dots, \underline{X}^{(n)}$ iid $\sim N(\underline{\mu}, \Sigma)$

Approximation to M , $\hat{M}_n = \sum_{i=1}^n I(X_{\max}^{(i)}/X_{\min}^{(i)} > c)/n$

Generalizes to more complicated joint distributions beyond $N_d(\underline{\mu}, \Sigma)$, provided we can simulate them.

And to more complicated functions (not just X_{\max}/X_{\min}).

Statistical Inference

Simple e.g. $Y_i = \beta X_i + \delta_i + \varepsilon_i$: $\delta_1, \dots, \delta_n \stackrel{iid}{\sim} \text{Gamma}(\lambda, \theta X_i)$
 $\varepsilon_1, \dots, \varepsilon_n \stackrel{iid}{\sim} N(0, \sigma^2)$

Observe: $(X_i, Y_i), i=1, \dots, n$

MLE for $\beta, \lambda, \theta, \sigma^2 = \underset{\beta, \lambda, \theta, \sigma^2}{\operatorname{arg\,max}} l(\beta, \lambda, \theta, \sigma^2; X, Y)$

Maximize via computer algorithm: e.g. Newton-Raphson alg.

Bayesian inf.: $\pi(\lambda, \beta, \theta, \sigma^2 | \dots) \propto l(\dots; (X_1, Y_1), \dots, (X_n, Y_n))$

MCMC to approximate $\pi(\dots | \dots) \propto p(\dots; \beta, \theta, \sigma^2) \leftarrow \text{prior distr.}$

OR optimization to maximize $\pi(\dots | \dots)$ w.r.t. $(\lambda, \beta, \theta, \sigma^2)$

Modern stats: lots of computing challenges.

Many topics I don't cover will be covered by you
in projects!

Projects: central to this course

Talk to me now about potential topics

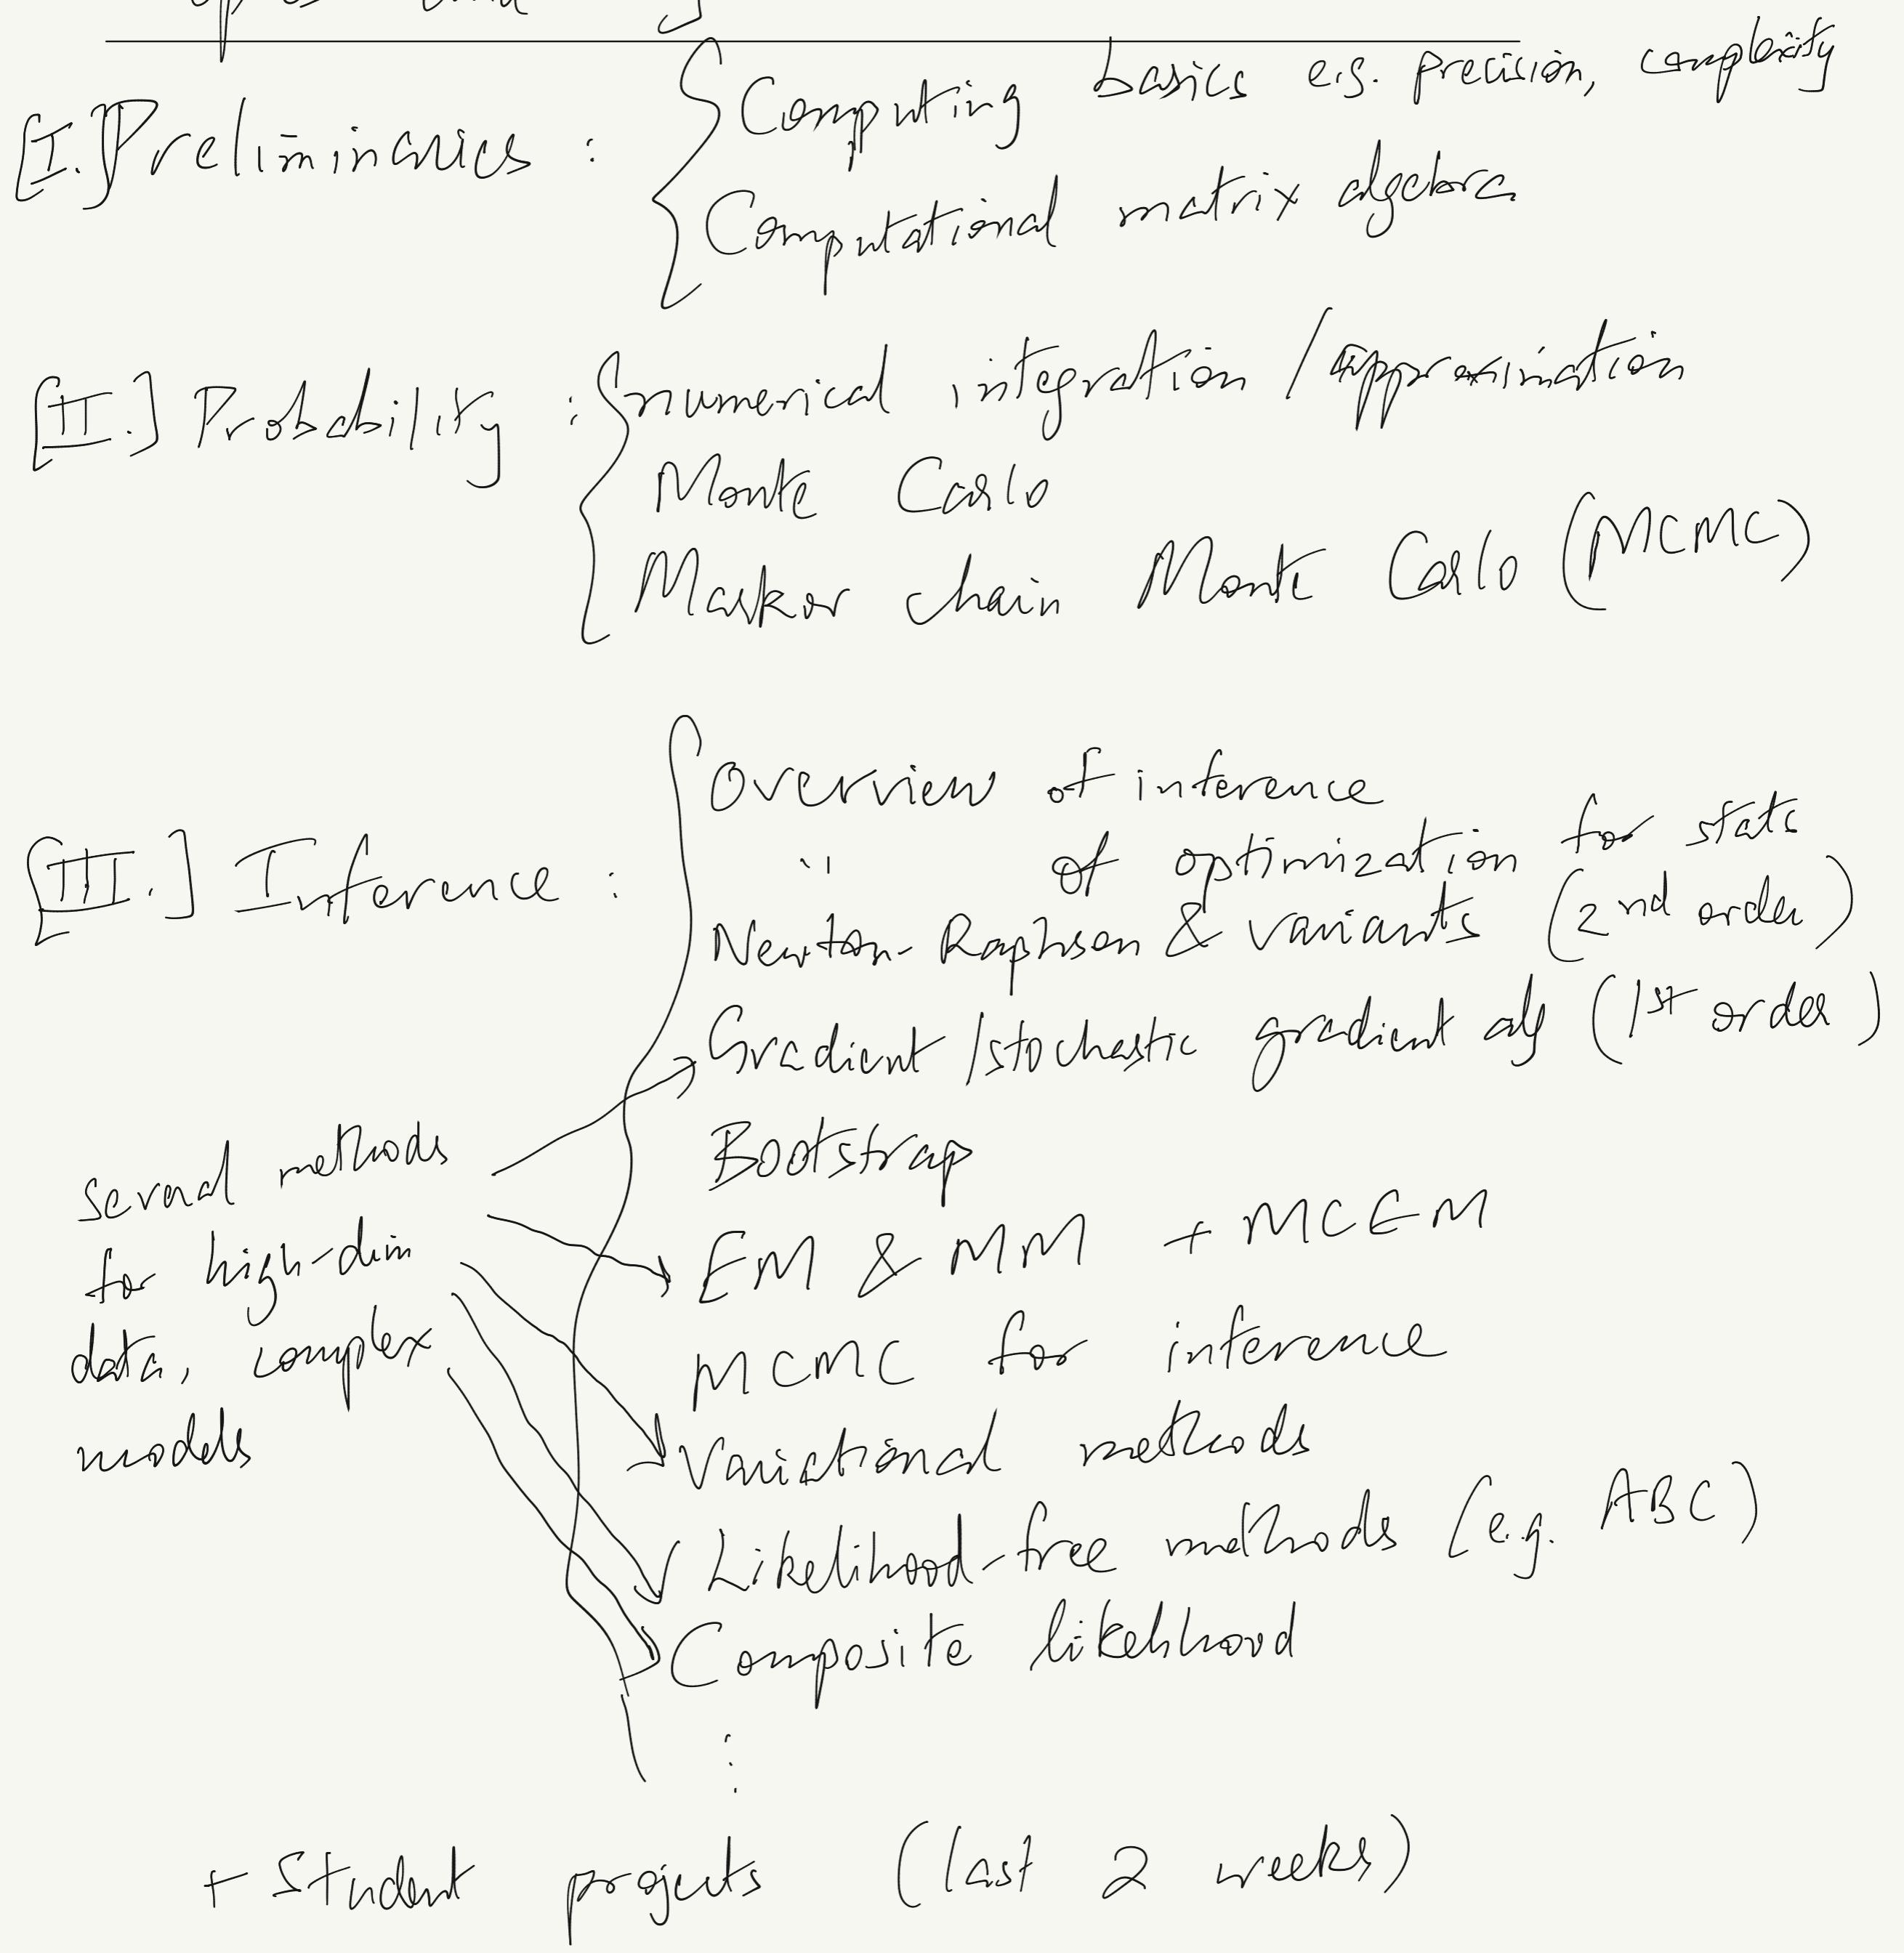
Substantial prelim. work: complete in 6 weeks

May have to modify topic so good

to start early. In past: poor projects typically
from students who do not discuss project w/ me.

Project must be primarily focused on algorithms —
nice applications are okay, but they can only be
motivation, not focus.

Topics and organization of course



Updates to schedule on course website

Computing preliminaries

- ① Computational complexity / cost of algorithms
 - Mathematical elegance is not same as computational efficiency
- ② Computer precision, stability
- ③ Pseudo-randomness
- ④ Computer memory
- ⑤ Parallel computing

End of Lecture 1
August 21, 2018

Clarity of thinking about computing issues vs statistical issues

"Separation of concerns": keeping statistical ideas separate from computing ideas (original quote from E.W. Dijkstra "A discipline of programming")
E.g. "My bootstrap CIs were narrower than my MCMC intervals"
Better: My percentile bootstrap CIs based on ML estimates were narrower than my 95% credible intervals from a Bayesian approach using MCMC.
E.g. I used EM to estimate parameters, then I used MCMC.
Better: I estimated params by MLG using the EM algorithm, then I used a Bayesian approach, implemented via MCMC.

Why: Same algorithm can be used for different purposes:

MCMC can be used for MLG too!

Also: clearer thinking separates computing issues e.g. Monte Carlo error from inference issue e.g. standard error.

① Computational complexity

Useful to use order notation for algorithm run times.

Let n be the unit of the # of operations

- "Big Oh" (\mathcal{O} -mega!) $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} < \infty$ $f: \mathbb{Z}^+ \rightarrow \mathbb{R}^+$
 $\mathcal{O}(g(n))$ is a class of functions $f(n) \in \mathcal{O}(g(n))$
 $g: \mathbb{Z}^+ \rightarrow \mathbb{R}^+$

for which $\exists k > 0, \exists n_0$ s.t. $\forall n > n_0,$
 $f(n) \leq k g(n)$

That is, f is bounded asymptotically from above by g .
 "grows at the same rate or slower than"

E.g. $n^3/3 \in \mathcal{O}(n^3)$
 $n^2 \in \mathcal{O}(n^2 + n)$
 $n^2 \in \mathcal{O}(10n^2)$

Big-Oh: "is of the same order as". Hence, popular for alg. cost
 $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 0$

- "Little oh" (\mathcal{o} -micron!)

$\mathcal{o}(g(n))$ is a class of functions $f(n) \in \mathcal{o}(g(n))$
 for which $\exists n_0$ s.t. $\forall n > n_0, \forall k > 0,$
 $f(n) \leq k g(n)$

That is, f is dominated asymptotically by g .

E.g. $n^2 \in \mathcal{o}(n^3)$
 $n^2 \in \mathcal{o}(n!)$
 $\ln(n) \in \mathcal{o}(n)$

Note: $n^2 \notin \mathcal{O}(n^3)$
 $n^2 \notin \mathcal{O}(n!)$
 $\ln(n) \notin \mathcal{O}(n)$

"grows strictly slower than"

Of course, $f \in \mathcal{o}(g) \Rightarrow f \in \mathcal{O}(g)$ but not vice-versa

- Theta notation

$\Theta(g(n))$ is a class of functions $f(n) \in \Theta(g(n))$ for which $\exists k_1 > k_2 > 0, \exists n_0$ s.t. $\forall n > n_0,$

$$k_2 g(n) \leq f(n) \leq k_1 g(n)$$

That is, f is bounded asymptotically from above and below by g .

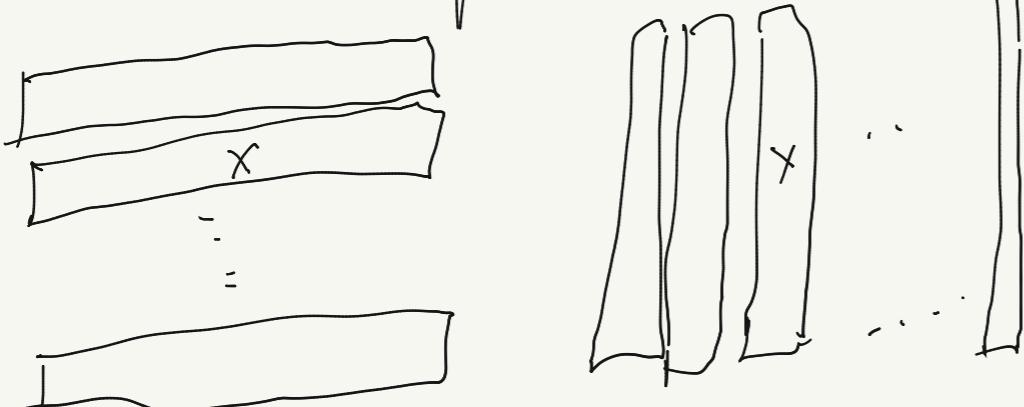
Why is this useful? Allows us to describe, in fairly simple terms, the computational cost of an algorithm. E.g. take two $n \times n$ matrices A, B .

$A+B$ or $A-B$ involves n^2 calculations

Each calculation is a floating point operation

Matrix multiplication : $C = A B$ Naive :

R matrix multiplication example
(flop)
1 flop each:
 $+,-,\times, /$ for 1R



$$C_{23} = [1 \ 2 \ \dots \ n]$$



= n multiplications +
 $(n-1)$ additions

Total # of calculations

$$= n^2 \times (2n - 1) \in O(n^3) \text{ flops}$$

So, costs increase at the rate n^3 w/ dimension of the matrix. Helps w/ understanding how alg. scales w/ size of matrix.

Common to use Strassen's alg. : $O(n^{2.8})$ but less stable and uses more memory - not good for very large matrices

What to do in practice about computing costs:

- ① Begin by "profiling" your code: this helps identify bottlenecks - parts of code that take the most time. Helpful since full-blown complexity analysis can be difficult.
In R: Rprof
- ② Be aware of computational costs of any potentially expensive parts of the code. E.g. Crokin and Van Loan book; Numerical recipes in C; Cormen, Leiserson, ... book
- ③ Counting flops not enough! (numerical solvers (e.g. BLAS, LAPACK) are optimized to hardware
 - (2) memory plays a role in costs (more on this later)
 - (3) file I/O can be expensive. e.g. pay attention to how often you write to file
- ④ algorithms can be iterative w/ some stopping criteria
e.g. eigen decomposition

- (4) If expensive operations are repeated, pre-calculate and re-use. E.g. some parts of likelihood may stay the same multiple iterations of an algorithm. $\ell(\theta, \phi; x) = \ell_1(\theta, \phi; x) + \ell_2(\theta; x) + \ell_3(\theta; x)$
- no need
to re-evaluate
if & same
- (5) Compiled language may often be much faster
E.g. C/C++/FORTRAN. Easy to do this in R.
Roger Peng's or C.J. Heyer's notes online to connect C code to R.
- General rule: your time is more valuable than computer time - code in compiled language if enough payoff!
- (6) Worth being aware of how languages handle matrices
- E.g. In R
$$\begin{bmatrix} & \\ & \end{bmatrix}_{1000 \times 1,000,000}$$
- is v. hard to read in but
$$\begin{bmatrix} & \\ & \end{bmatrix}_{1000,000 \times 1,000}$$
 is manageable

End of lecture 2,
August 23, 2018

② Computer precision / stability

Computers are not infinitely precise. This can affect accuracy of calculations.

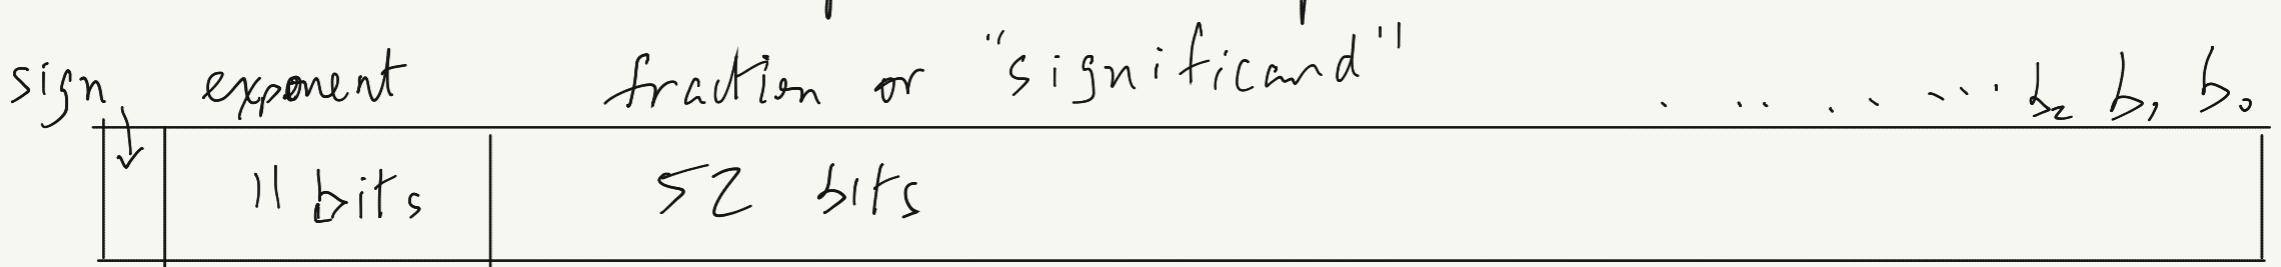
Basics: represent #s as a sequence of 0's and 1's
 "bits": term coined by statistician John Tukey

\downarrow
 binary digits

[Group of 8 bits is a "byte" (typically) — historically used to represent a single text character]

Representation of #s either 32 bit or 64 bit

FEEE-754 double precision representation:



$$(-1)^{\text{sign}} \left[1 + \sum_{i=1}^{52} b_{52-i} 2^{-i} \right] \times 2^{\text{expo} - 1023}$$

E.g. $\lfloor 1000\ldots01 \rfloor \quad \overbrace{00\ldots1000}^1 \rfloor$

$$(-1)^1 \times \left[1 + \left(\frac{1}{2}\right)^{48} + \left(\frac{1}{2}\right)^{52} \right] \times 2^{(1024+4+1)-1023}$$

$$\left[= -1.0\ldots037746 \times 2^6 \right]$$

E.g. Exponent $0\ldots1 : 2^{-1022}$
 $11\ldots10 : 2^{2046-1023} = 2^{1023}$

E.g. $0 \underbrace{01\ldots11}_{1023-1023} \underbrace{0\ldots-010}_{1+0+\dots+(\frac{1}{2})^{51}+0} = 1.0\ldots04\ldots$
 $2 \times \left[1 + 0 + \dots + \left(\frac{1}{2}\right)^{51} + 0 \right] \approx 1 + 4 \times 10^{-16}$

In practice : need to pay attention to

(1) how precisely numbers are stored - this can affect results.

(2) how numbers are displayed. In R:

good to learn how to use: paste, format, sprintf
- default presentation of numbers can obscure what is really going on.

Add R examples

(3) common to encounter numbers too large

or small to be represented on a computer,

Good to work in log scale in many problems

e.g. frequently in likelihood calculations for complicated models.

$$\frac{L(\theta; \underline{x})}{L(\phi; \underline{x})},$$

use $\ell(\theta; \underline{x}) - \ell(\phi; \underline{x})$

Stability

Calculation carried out on a computer will, due to rounding error, produce results that are different from results based on theory.

Of interest to us: how close is the computer result to the theoretical result?

[Reference: Stewart & Voss, Numerical Linear Algebra]

Backward error analysis: view computed solution to a problem as exact solution to a perturbed problem.

Small perturbation \Rightarrow algorithm is "backward stable"
Not small perturbation \Rightarrow algorithm is "unstable"

We will discuss this more in context of numerical linear algebra.

③ Pseudo-randomness

Deterministic algorithms are used to mimic random sequences on a computer

Algorithms look roughly like this:

- ① Start w/ initial value u_0 , called the "seed". Like an initial state. Vector of integers
- ② Construct a deterministic sequence u_1, u_2, \dots where successive values use previous values

Simple e.g. $u_{k+1} = f(u_k)$, $k = 1, 2, \dots$

- ③ Resulting sequence of integers $\in [0, u_{\max}]$ where u_{\max} is set by the algorithm, typically related to how numbers are represented

$$\subset [0, 1]$$

Sequence: $u_0/u_{\max}, u_1/u_{\max}, \dots$

Good random number generator: sequence is virtually

"indistinguishable" from $\text{Unif}(0,1)$ iid sequence.

\Leftrightarrow Pass several tests e.g. "Diehard battery" (Marsaglia, 1960s, 1970s)

Very popular: Mersenne-Twister (1997) default in R, Python, Matlab . . .

All random number generation algorithms are based on iid sequence of $\text{Unif}(0,1)$ random variables

Practical consequences:

- ① Deterministic: if use same seed, get same sequence
But this is an advantage more than a
disadvantage: helps w/ reproducible work.
Also, useful for debugging code.

In R: `set.seed(x)`

• `Random.seed` is global variable that
stores state of random number generator

Very useful for long, potential 'interrupted' runs
of MCMC, or stochastic gradient alg.

Care w/ parallel computing: make sure different seeds on
different processors!

- ② Random number generator has a "period": eventually
returns to 'initial state', then reproduces same
sequence: generally a non-issue

E.g. Mersenne-Twister's period is 2^{19937}

End of Lecture 3,
Tuesday August 28,
2018

④ Computer memory basics

Hierarchy : cache : small, v. fast access } O.S. keeps most frequently accessed items here
RAM : bigger, next fastest } next, O.S. uses RAM

Hard drive / swap space: When oper. sys. runs out of RAM,
internal hard drive is used to store items

temporarily not being used. Writing / retrieving items
from various kinds of memory is called **swapping**.

Practical consequences:

Useful to know if you are using a lot of RAM, e.g. by
constantly accessing huge matrices. This can result in a
lot of swapping, which can really slow things down.

This is an e.g. of why calculating real computational
costs of an algorithm can be very complicated.

- Careful of algorithms that require lots of memory
E. Strassen's alg vs simple matrix multiplication.

Often tradeoff between memory use and flops.

- Avoid too much dynamic allocation of memory.

E.g. in R

```
myvec = c(1)
for (i in 2:m) {
  myvec = c(myvec, myvec[i-1] + i)
```

dynamic memory allocation

Better :

```
myvec = rep(NA, m) # fixed allocation of memory
```

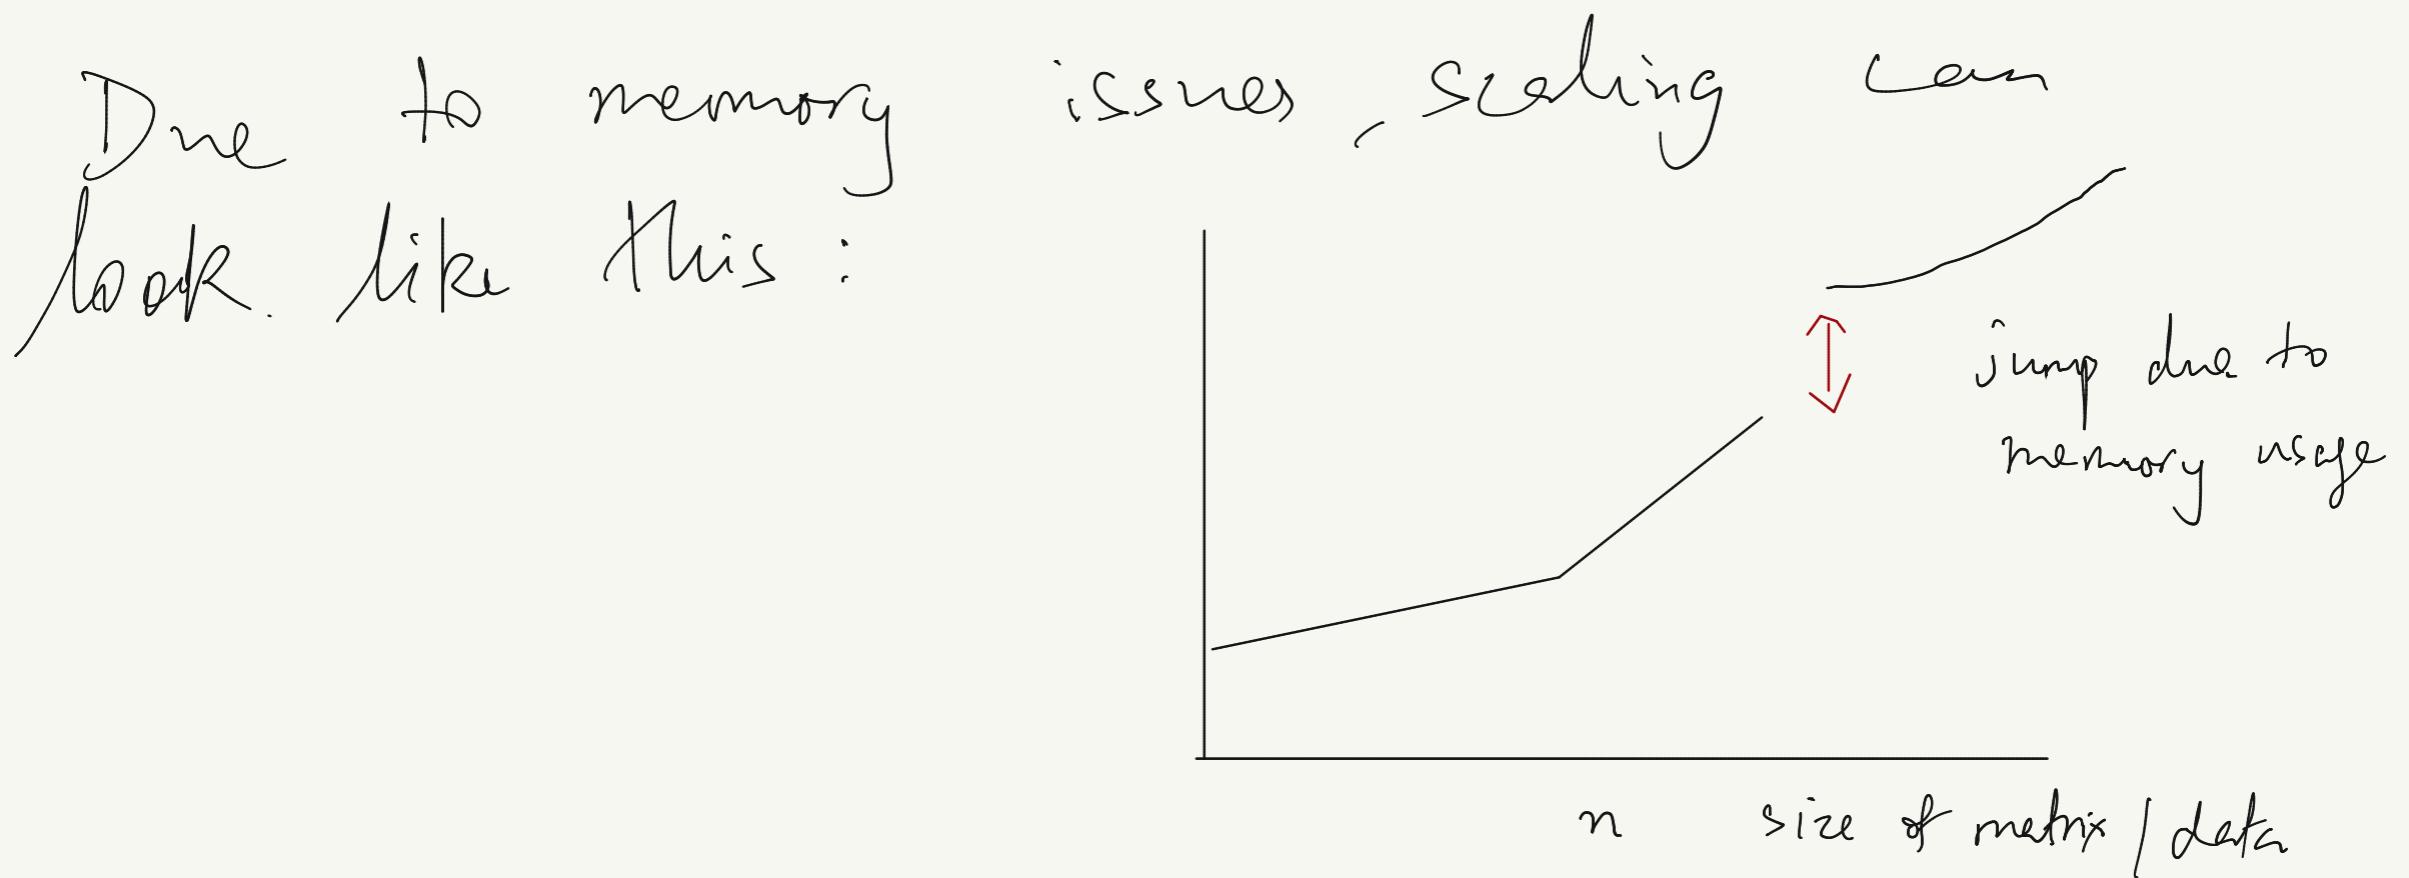
```
myvec[1] = 1
```

```
for (i in 2:m)
```

```
  myvec[i] = myvec[i-1] + i
```

In practice, common to not know length/size of data structures in advance. Still good to avoid many memory allocations: allocate estimated size in advance, then increase if necessary, reduce how often this is done.

Point process problem??



Parallel computing

High performance computing (HPC) broadly refers to the use of multiple processors to solve challenging computing problems.

Learning about how to parallelize algorithms effectively is therefore a very useful skill in modern scientific computing.

In statistics, lots of opportunities for parallel computing.

Embarrassingly parallel problem: one where very little effort is required to separate the problem into a number of parallel tasks.

E.g. matrix multiplication $O(nm k^2)$ flops

$$\begin{bmatrix} C & \dots & \dots \\ \vdots & \ddots & \vdots \\ \vdots & \vdots & \vdots \end{bmatrix}_{n \times m} \xrightarrow{n \times k} A_{n \times k} \quad \left\{ \begin{bmatrix} \vdots \\ \vdots \\ \vdots \\ \vdots \end{bmatrix}_{k \times m} \right. \quad = \quad C_{n \times m}$$
$$B_{k \times m}$$

Each element, C_{nm} , in matrix, may be computed independently. Easy to parallelise. If $n m$ processors: $k^2 +$ overhead for distributing and gathering results

Cost depends on # of processors
 Speed / capacity of processors
 requirements of parallelized computations
 communication overhead

Gains depend on size of problem + available computing
 (Andahl's bound, 1967)

$$\text{E.g. if } L(\theta; \mathbf{x}) = \prod_{i=1}^n L(\theta_i; x_i)$$

$L(\theta; \mathbf{x}) = \sum_{i=1}^n L(\theta_i; x_i) \leftarrow$ parallelize these calculations,
 each of which might be expensive

E.g. Simulation studies : see RMSE of estimate $\hat{\theta}$ as

true value of θ varies

$$\begin{array}{l} \theta_1 \rightarrow \text{RMSE}(\theta_1) \\ \vdots \\ \theta_K \rightarrow \text{RMSE}(\theta_K) \end{array} \quad \left. \begin{array}{l} \text{each on a different} \\ \text{processor} \end{array} \right\}$$

Matlab & R have capabilities for simple parallelization
 e.g. snow (simple network of workstations) package.

Graphical processing units (GPUs) : huge number of processors,
 each w/ very limited capabilities

When not embarrassingly parallel, can require a
 lot more creativity

Numerical Linear Algebra

Let real matrix $A \in \mathbb{R}^{m \times n}$ w/ $m \geq n$
 $b \in \mathbb{R}^m$

Three main problems:

① Solving system of linear equations. Let $m = n$

Find $x \in \mathbb{R}^n$ s.t. $Ax = b$

E.g. matrix inversion in multivariate normal pdf

$$f(x; \mu, \Sigma) = (2\pi)^{-n/2} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} (x - \mu)^\top \Sigma^{-1} (x - \mu) \right\}$$

In practice, should try to avoid taking inverses

② Least Squares.

Find $x \in \mathbb{R}^n = \arg \min_{x \in \mathbb{R}^n} \|Ax - b\|^2$

E.g. regression notation: find $\hat{\beta} = \arg \min_{\beta \in \mathbb{R}^d} \|y - X\beta\|^2$

Where X : $n \times p$ matrix, p predictors, n data points

β : p -dimensional vector of reg. coefficients

y : n -dimensional " " responses

In this case, simple calculus yields: $\hat{\beta} = (X^\top X)^{-1} X^\top y$

So inversion again. Very common in high-dimensional problems to have numerical issues when computing

$(X^\top X)^{-1}$, $X^\top X$ is often "ill-conditioned"

If $p > n$: stat issue (non-invertible + not enough data to relate p predictors)

For large n, p : may be invertible in theory but numerical issues

(3) Eigenvalue problem: Let $m = n$
Find $(x, \lambda) \in \mathbb{R}^n \times \mathbb{R}$
 $Ax = \lambda x, \|x\|_2^2 = 1$

Eigen decompositions are central to principal component analysis.

E.g. dimension reduction via principal component analysis
If # of predictors is large, and many are correlated,
use small # of linear combinations of variables
that captures most of the variability of the data.

In statistics it is very common to have structured matrices. E.g. covariance matrices may be patterned.

Learn to take advantage of structures to speed up computation.

E.g. of methods for taking advantage of structured matrices.

Spatial or variance components model

$$(x_1, \dots, x_n)^T \sim N(\mu, \Sigma)$$

where $\Sigma_{n \times n} = \sigma^2 I + H$, where

H can be written as $\underset{n \times k}{K} \underset{k \times n}{K^T}$ (This is common in spatial stats models)

To evaluate pdf, calculate Σ^{-1}

Cost: $O(n^3)$

Faster: use Sherman-Morrison-Woodbury formula

$A_{n \times n}, C_{k \times k}$: invertible

Let $U : n \times k$ matrix

$V : k \times n$ matrix

$$(A + UCV)^{-1} = A^{-1} - A^{-1}U(C^{-1} + VA^{-1}U)^{-1}VA^{-1}$$

\uparrow
k × k matrix

Very useful if A^{-1} is easy e.g. A is diagonal matrix
and $R \ll n$

End of Thursday,
Aug 30, 2018
lecture

Note: Sylvester's formula: determinant for same form of matrix.

$$\text{Above e.g. } \Sigma^{-1} = (\sigma^2 I)^{-1} - (\sigma^2 I)^{-1} K \underbrace{\left(I^{-1} + K^T (\sigma^2 I)^{-1} K \right)^{-1}}_{\text{inversion of j × j matrix}} K (\sigma^2 I)^{-1}$$

Dominant cost: inversion of $j \times j$ and
matrix multiplications

Q. Calculate exact cost?

Partitioned matrices

Let A be invertible matrix, partitioned as follows

$$A_{n \times n} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}$$

E.g. when adding new rows (observations) in linear repr.
or " columns (variables) in "

Rather than work w/ full A , use structure of matrix
to get

$$A^{-1} = \begin{bmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{bmatrix} \text{ where}$$

$$C_{11} = (A_{11} - A_{12} A_{22}^{-1} A_{21})^{-1}$$

$$C_{22} = (A_{22} - A_{21} A_{11}^{-1} A_{12})^{-1}$$

$$C_{12} = -A_{11}^{-1} A_{12} C_{22}$$

$$C_{21} = -A_{22}^{-1} A_{21} C_{11}$$

Note:
Sherman-
Woodbury-
Morrison
is a special case where $A^* = \begin{bmatrix} A & u \\ v & -c^{-1} \end{bmatrix}$

Matrix factorizations (stats-centric view)

Cholesky factorization: If $A \in \mathbb{R}^{n \times n}$ is positive definite, then there exists a unique lower triangular $C \in \mathbb{R}^{n \times n}$ w/ positive diagonal elements s.t.

$$A = CC^T \quad (\text{see Stuart \& Voss for proof})$$

C is a matrix "square root"

Some uses of Cholesky factorization:

[1] Simulating $\tilde{x} \sim N_n(\mu, \Sigma)$ $\mu \in \mathbb{R}^n$, $\Sigma \in \mathbb{R}^{n \times n}$

(1) Assume we can simulate $z_1, \dots, z_n \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$ $\Sigma \succ 0$

$$\underline{z} = (z_1, \dots, z_n)^T \sim N(0, I)$$

(2) Find Cholesky factor of Σ , C s.t. $CC^T = \Sigma$

$$(3) C\underline{z} \sim N(0, CC^T)$$

$$\text{So } C\underline{z} \sim N(0, \Sigma)$$

$$(4) \underline{x} = C\underline{z} + \mu \sim N(\mu, \Sigma)$$

Computational cost of above algorithm: dominated by

Cholesky decomposition $O(n^3)$.

Related: Simulate $\underline{x} \sim N(0, Q^{-1})$ where Q is a "precision matrix". Common in Markov random field models

Find C s.t. $CC^T = Q$. Solve: $Cy = z$, \leftarrow easy because C is lower triangular

where $\underline{z} \sim N(0, I)$. Clearly, $y \sim N(0, Q^{-1})$

[2] Evaluating multivariate normal pdfs

$$f(x; \mu, \Sigma) = (2\pi)^{-n/2} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} (x-\mu)^T \Sigma^{-1} (x-\mu) \right\}$$

(1) Find Cholesky factor of Σ : C

$$\begin{aligned} (2) \text{ Now, } (x-\mu)^T \Sigma^{-1} (x-\mu) \\ &= (x-\mu)^T (CC^T)^{-1} (x-\mu) \\ &= (x-\mu)^T (C^T)^{-1} C^{-1} (x-\mu) \end{aligned}$$

So, solve $Cx = (x-\mu)$ for x

$$\text{Then } x = C^{-1}(x-\mu)$$

Compute $x^T x$, which is $[C^{-1}(x-\mu)]^T [C^{-1}(x-\mu)]$

$$\begin{aligned} &= (x-\mu)^T [C^{-1}]^T C^{-1} (x-\mu) \\ &= (x-\mu)^T \Sigma^{-1} (x-\mu) \end{aligned}$$

(3) To find $|\Sigma|^{-1/2}$, note that $|\Sigma|^{-1/2} = |CC^T|^{-1/2}$

where C was already computed above

$$\text{Now } |CC^T|^{-1/2} = |C|^{-1}$$

And $|C|$ is just a product of its diagonal elements since it is triangular.

Main computational cost is due to Choleski factor again, $O(n^3)$.

[3]

Avoiding matrix inversion

Want x s.t. $Ax = b$

That is, $x = A^{-1}b$

(1) write $A = L U$ Choleski factor

(2) solve $L y = b$ for y ($y = L^{-1}b$)

(3) solve $U x = y$ for x
 $(x = U^{-1}y = U^{-1}L^{-1}b = A^{-1}b)$

Bank to examples of structured matrices

Already discussed partitioned matrices and
matrices amenable to Sherman-Woodbury-Morrison

identify → faster inversion and determinants

Now: structured matrices for speeding up Cholesky
factorization.

E.g. banded matrices

Consider AR(1) process, a 1st order Markov chain

$x_1 \sim N(0, \sigma^2)$, $x_n | x_{n-1} \sim N(\phi x_{n-1}, \sigma^2)$ for $n=2, 3, \dots$
w/ $\phi \in (-1, 1)$, fixed constant

$\Sigma = \text{Cor}(x_1, x_2, \dots, x_n)$ is a "dense" matrix.

That is, $\Sigma = \begin{bmatrix} \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot \end{bmatrix}$
 $\Sigma_{ij} > 0 \quad \forall i, j$

Interestingly, precision matrix $Q = \Sigma^{-1}$
is sparse (lots of zeroes). In fact it
is a "banded" matrix.

$$Q = \begin{bmatrix} x & x & & & \\ x & \ddots & & & \\ & \ddots & \ddots & & \\ & & \ddots & \ddots & x \\ & & & x & x \end{bmatrix} \quad \text{band}$$

This leads itself to much faster
Choleski factorization algorithms, $O(n b_w^2)$
where $b_w = \text{bandwidth}$. Above $b_w = 1$

E.g. if $n = 5,000$
 $n^3 = 1.25 \times 10^9$

$n b_w^2 = 5,000$ even if $b_w = 10$, $n b_w^2 = 5 \times 10^5$

Enormous storage reductions as well.

For $b_w = 1$ only need to store $\leq 3n$

Above was a simple example but sparse and banded matrices are common in many fields

e.g. graphical models, spatial models.

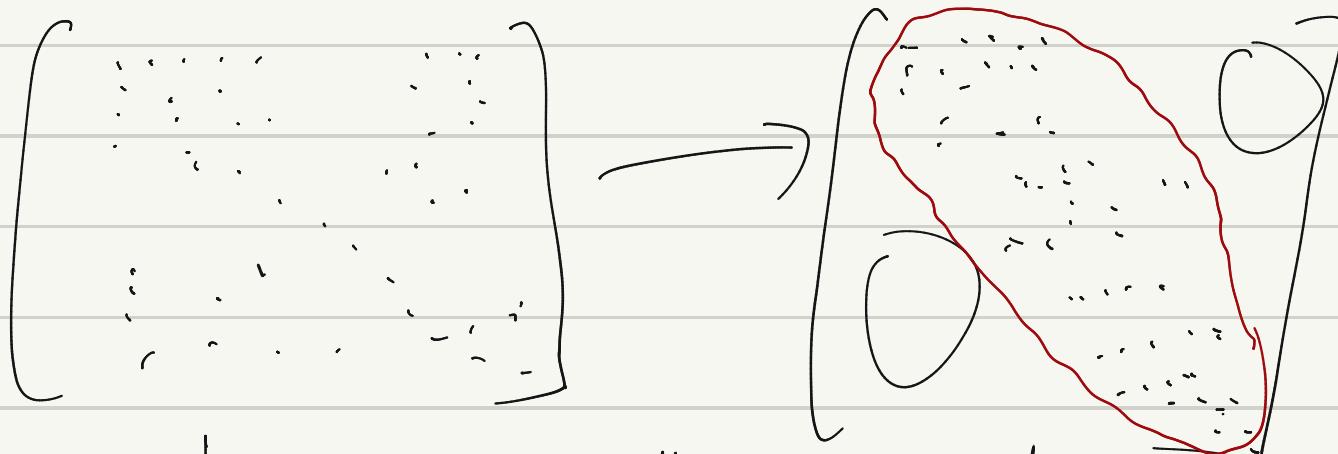
E.g. models where dependence is characterized mostly by conditional independence structure

Even if matrix is not banded, it might be sparse. Possible to reduce bandwidth

using an algorithm e.g.

Gibbs-Poole-Stockmeyer-King / Cuthill-McKee

Fixed one-time cost



b_w may remain small even as n gets large.

Stack overflow: invert a matrix w/ reduced bandwidth
"Gibbs Poole Stockmeyer alg in R"

Popular structure in statistics: Kronecker product for covariances.

$A_{m \times n}$, $B_{p \times q}$, then Kronecker product $A \otimes B$ is
 $m p \times n q$ block matrix

$$C = A \otimes B = \begin{bmatrix} a_{11}B & \dots & a_{1n}B \\ \vdots & & \vdots \\ a_{m1}B & \dots & a_{mn}B \end{bmatrix}_{mp \times nq}$$

Matrix operations are much faster

$$C^{-1} = \underbrace{B^{-1} \otimes A^{-1}}_{\text{dominated by cost of larger of } B, A \text{ invertible}}, \text{ if } C \text{ is invertible} (\Leftrightarrow B, A \text{ invertible})$$

e.g. $A_{m \times m}$, $B_{g \times g}$, $m > g$, then cost in $B^{-1} \otimes A^{-1}$
 is dominated by m^3 . Cost of direct C^{-1} is order
 $(mg)^3$ which may be much larger than m^3

$$|C| = |A|^m |B|^n \text{ if } A_{n \times n}, B_{m \times m}$$

Computational advantages + useful modeling framework

e.g. A describes one kind of dependence, while

B .. another ,

$A \otimes B$ combines both

Eigendecomposition

Most popular use in stats is perhaps principal components analysis (PCA)

Let X have covariance matrix Σ

$\therefore \Sigma$ is symmetric, can be decomposed as follows

$$\Sigma = U D U^T$$

where D is diagonal matrix of eigenvalues of Σ and U is corresponding orthogonal matrix of eigenvectors

Suppose eigenvalues are distinct and ordered

$$0 \leq \lambda_1 < \lambda_2 \dots < \lambda_m \text{ then}$$

columns u_1, \dots, u_m are unique up to sign (cf. k-Large book)

$$\text{R.v.s. } V_j = u_j^T X, j=1, \dots, m$$

hence covariance matrix $U^T \Sigma U = D$

V_1, \dots, V_m are principal components

uncorrelated, and increasing in variance.

- HW exercises:
- (1) Return package dti data set
 - (2) eigenvalue distr.
 - (3) Sim study
 - (4) Kernel rep : use of Sherman-Morrison formula