

# Proximal Methods

---

Tobia Boschi

Stat 540 - Spring 2108

# Why Proximal Methods?

- ⊙ **High dimensional convex problems**
  - non-differentiable
  - constrained
  - large-size and parallel implementations
- ⊙ **Proximal methods to solve LASSO**

## Key Idea

- avoid *gradient* and *hessian* computation
  - evaluate instead the *proximal operator*:
- small convex optimization problem

## Definition

Let  $f : R^n \rightarrow R$  be a closed proper *convex* function.

The **proximal operator**  $\text{prox}_{\lambda f} : R^n \rightarrow R^n$  is defined by:

$$\text{prox}_{\lambda f}(x) = \operatorname{argmin}_z \left( f(z) + \frac{1}{2\lambda} \|z - x\|_2^2 \right)$$

It balances two goals:


- ⊙ minimizing  $f$
- ⊙ staying near  $x$

### Possible interpretations

- surrogate method:  $f(x) \longrightarrow f(z) + \frac{1}{2\lambda} \|z - x\|_2^2$
- gradient step for  $f$ :  $\text{prox}_{\lambda f}(x) \cong x - \lambda \nabla f(x)$

# Simple Example

- ⊙  $f(x) = |x|$
- ⊙  $|z| + \frac{1}{2\lambda}(x - z)^2$
- ⊙  $\text{prox}_{\lambda f}(x) = \text{sign}(x)(|x| - \lambda)_+$




# Simple Example

⊙  $f(x) = |x|$


⊙  $|z| + \frac{1}{2\lambda}(x - z)^2$

⊙  $\text{prox}_{\lambda f}(x) = \text{sign}(x)(|x| - \lambda)_+$




# Simple Example

- ⊙  $f(x) = |x|$
- ⊙  $|z| + \frac{1}{2\lambda}(x - z)^2$
- ⊙  $\text{prox}_{\lambda f}(x) = \text{sign}(x)(|x| - \lambda)_+$



# Simple Example

- ⊙  $f(x) = |x|$
- ⊙  $|z| + \frac{1}{2\lambda}(x - z)^2$
- ⊙  $\text{prox}_{\lambda f}(x) = \text{sign}(x)(|x| - \lambda)_+$



# Proximal Methods and Regression

Let  $x \in R^n, b \in R^m, A \in R^{m \times n}$ . We want to minimize:

$$g(x) + h(x) = \frac{1}{2} \|Ax - b\|_2^2 + h(x)$$

- ⊙ LASSO:  $h(x) = \gamma \|x\|_1$
- ⊙ RIDGE:  $h(x) = \frac{\gamma}{2} \|x\|_2^2$
- ⊙ ELASTIC:  $h(x) = \gamma_1 \|x\|_1 + \frac{\gamma_2}{2} \|x\|_2^2$

## Implemented Algorithms

- Proximal Gradient
- Proximal ADMM  
*(alternating direction method of multipliers)*



# Proximal Methods and Regression

Let  $x \in R^n, b \in R^m, A \in R^{m \times n}$ . We want to minimize:

$$g(x) + h(x) = \frac{1}{2} \|Ax - b\|_2^2 + h(x)$$

- ⊙ LASSO:  $h(x) = \gamma \|x\|_1$
- ⊙ RIDGE:  $h(x) = \frac{\gamma}{2} \|x\|_2^2$
- ⊙ ELASTIC:  $h(x) = \gamma_1 \|x\|_1 + \frac{\gamma_2}{2} \|x\|_2^2$

## Implemented Algorithms

- Proximal Gradient
- Proximal ADMM  
(*alternating direction method of multipliers*)

# Proximal Gradient

## Pseudo-code

- ⊙ Gradient step

$$v^k = x^k - \lambda \nabla g(x^k)$$

- ⊙ Proximal operator step

$$x^{k+1} = \text{prox}_{\lambda h}(v^k)$$

# Proximal Gradient

## Pseudo-code

- ⊙ Gradient step

$$v^k = x^k - \lambda \nabla g(x^k)$$

- ⊙ Proximal operator step

$$x^{k+1} = \text{prox}_{\lambda h}(v^k)$$

- $\nabla g(x) = A'Ax - A'b$ 
  - precompute  $A'A$  and  $A'b$
  - at each iteration evaluate  $A'Ax$ :  $\mathcal{O}(n^2)$
- $\left( \text{prox}_{\lambda \gamma \|\cdot\|_1}(x) \right)_i = \text{prox}_{\lambda \gamma |x_i|}(x) = \text{sign}(x)(|x| - \lambda \gamma)_+$

# Proximal Gradient

- ⊙ Proximal Gradient as an **Majorization-Minimization** algorithm

$$\hat{g}_\lambda(x, y) = g(y) + \nabla g(y)'(x - y) + \frac{1}{2\lambda} \|x - y\|_2^2 \geq g(x)$$

*Majorization step*

- compute  $\hat{g}_\lambda(x, x^k) + h(x)$

*Minimization step*

- $\min_x \left( \hat{g}_\lambda(x, x^k) + h(x) \right) = \text{prox}_{\lambda h} \left( x^k - \lambda \nabla g(x^k) \right)$   
 $= \text{prox}_{\lambda h}(v^k)$

# Proximal Gradient

- ⊙ Proximal Gradient as an **Majorization-Minimization** algorithm

$$\hat{g}_\lambda(x, y) = g(y) + \nabla g(y)'(x - y) + \frac{1}{2\lambda} \|x - y\|_2^2 \geq g(x)$$

## *Majorization step*

- compute  $\hat{g}_\lambda(x, x^k) + h(x)$

## *Minimization step*

- $\min_x \left( \hat{g}_\lambda(x, x^k) + h(x) \right) = \text{prox}_{\lambda h} \left( x^k - \lambda \nabla g(x^k) \right)$   
 $= \text{prox}_{\lambda h}(v^k)$

- ⊙ **Backtracking of  $\lambda$**

- $x = \text{prox}_{\lambda h} (x^k - \lambda \nabla g(x^k))$
- *reduce  $\lambda$  since:*  $g(x) \leq \hat{g}_\lambda(x, x^k)$

## Proximal ADMM

**Note:** minimize  $g(x) + h(x)$  is equivalent to minimize:

$$g(x) + h(z) \quad \text{subject to} \quad x - z = 0$$

# Proximal ADMM

**Note:** minimize  $g(x) + h(x)$  is equivalent to minimize:

$$g(x) + h(z) \quad \text{subject to} \quad x - z = 0$$

## Pseudo-code

- ⊙  $x^{k+1} = \text{prox}_{\lambda \mathbf{g}}(z^k - u^k)$
- ⊙  $z^{k+1} = \text{prox}_{\lambda \mathbf{h}}(x^{k+1} + u^k)$
- ⊙  $u^{k+1} = u^k + x^{k+1} - z^{k+1}$

- $g(x) = \frac{1}{2} \|Ax - b\|_2^2 = \frac{1}{2} (b - Ax)'(b - Ax)$
- $\text{prox}_{\lambda \mathbf{g}}(x) = (I_n + \lambda A'A)^{-1}(x + \lambda A'b)$

# Proximal ADMM

How to compute  $(I_n + \lambda A'A)^{-1}$ :

⊙  $n > m$

- $C = \text{chol}(I_n + \lambda A'A)$
- $(I_n + \lambda A'A)^{-1} = (C')^{-1}C^{-1}$
- $\mathbb{O}(n^3)$

⊙  $m > n$

- inversion lemma:  $(I_n + \lambda A'A)^{-1} = A'(I_m + \lambda AA')^{-1}A$
- $C = \text{chol}(I_m + \lambda AA')$
- $\mathbb{O}(m^3)$



## Simulation 1

- $A = \text{random normal}(m \times n)$
- $x_0 = \text{random normal}(n \times 1)$
- $b = Ax_0 + \text{err}$
- sparsity = 0.95
- $\lambda = 1$
- $\gamma = 0.1 \|A'b\|_\infty$

⊙  $m = 500, n \uparrow$

	sec			obj		
	CVX	grad	ADMM	CVX	err grad	err ADMM
$n = 10^2$	0.10966	0.00136	0.02259	0.4606	0	0
$n = 10^3$	11.10568	0.02289	0.02589	8.5773	0.0010	0.0003
$n = 10^4$	155.49445	3.97270	0.57923	86.8624	0.1288	0.00665
$n = 4 \cdot 10^4$	820.12007	111.32019	4.21288	305.0444	1.0231	0.08624

## Simulation 1

- $A = \text{random normal}(m \times n)$
- $x_0 = \text{random normal}(n \times 1)$
- $b = Ax_0 + \text{err}$
- sparsity = 0.95
- $\lambda = 1$
- $\gamma = 0.1 \|A'b\|_\infty$


⊙  $m = 500, n \uparrow$

	sec			obj		
	CVX	grad	ADMM	CVX	err grad	err ADMM
$n = 10^2$	0.10966	0.00136	0.02259	0.4606	0	0
$n = 10^3$	11.10568	0.02289	0.02589	8.5773	0.0010	0.0003
$n = 10^4$	155.49445	3.97270	0.57923	86.8624	0.1288	0.00665
$n = 4 \cdot 10^4$	820.12007	111.32019	4.21288	305.0444	1.0231	0.08624

$n \gg m \Rightarrow \text{ADMM } \mathcal{O}(m^3) > \text{Gradient } \mathcal{O}(n^2)$

# Distributed ADMM


⊙ reduce  $m$   $\longrightarrow$  divide  $A$  in  $S$  blocks



- $g(x) = \sum g_i(x)$   
 $= \sum \frac{1}{2} \|A_i x - b_i\|_2^2$
- $\mathcal{O}((m/S)^3)$

# Distributed ADMM

⊙ reduce  $m \rightarrow$  divide  $A$  in  $S$  blocks



- $g(x) = \sum g_i(x)$   
 $= \sum \frac{1}{2} \|A_i x - b_i\|_2^2$
- $\mathcal{O}((m/S)^3)$

## Pseudo-code

- ⊙  $x_i^{k+1} = \text{prox}_{\lambda g_i}(z^k - u_i^k)$
- ⊙  $z^{k+1} = \text{prox}_{\frac{\lambda h}{S}}(\bar{x}^{k+1} + \bar{u}^k)$
- ⊙  $u_i^{k+1} = u_i^k + x_i^{k+1} - z^{k+1}$

## Simulation 2

⊙  $m = 10.000$ ,  $n = 50.000$

⊙  $S = 10$

sec		
grad	ADMM	distr
3463	423	90

obj		
grad	ADMM	distr
609.012	607.046	609.747

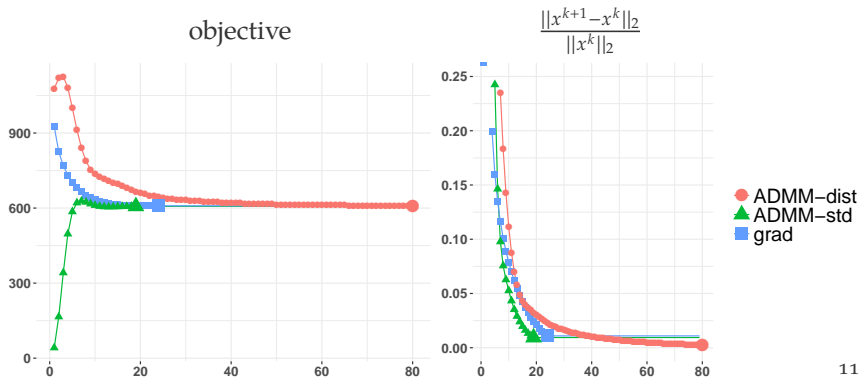
## Simulation 2

⊙  $m = 10.000$ ,  $n = 50.000$

⊙  $S = 10$

sec		
grad	ADMM	distr
3463	423	90

obj		
grad	ADMM	distr
609.012	607.046	609.747



# Conclusions

## Pros

- non-smooth problem
- much faster
- easy to parallelize
- good approximations

## Cons

- convex problems
- pointwise estimation
- approximations

## More work...

- sensitivity study for  $\lambda$  and  $\gamma$
- add performance indicators (*prediction error*)
- study constrained problems (*matrix decomposition*)