
MCMC strategies for Task-Specific Region Partition Problem

— Hua Wei —
College of IST, Penn State University

Learning Task-Specific City Region Partition

Task: Crime Prediction

Given:


- **tract** features and crime number in 2010
- **tract** features in 2011
- **Communities** consist of **tracts**

Task:

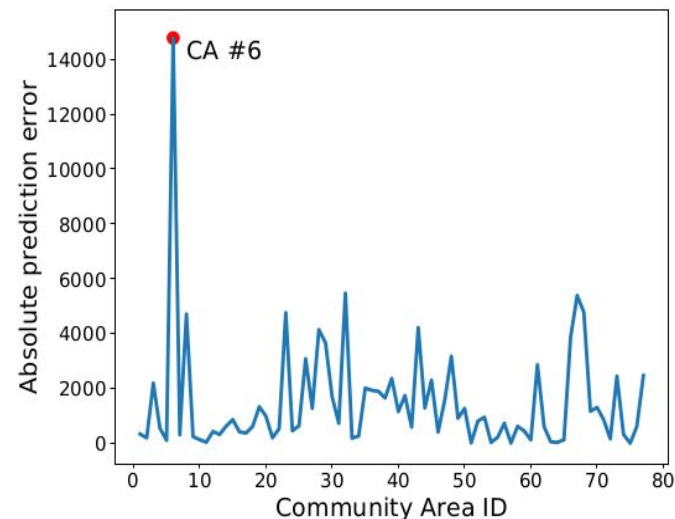
predict crime number for every **community** in 2011

Easy solution:

- Aggregate **tract-level** data into **community-level** features and crime number as training data
- Fit regression model



Left: crime prediction error at community level.



Right: community #6 is an outlier.


Learning Task-Specific City Region Partition

Explain the crime prediction error.


The eastside of Community #6 is different from others.

Question:

How do we obtain appropriate partitions?



(a) Census tracts of community #6.



(b) Tract similarity.

Appropriate partitions are needed!

- **Misalignment problem:** different observations are not in the same scale.
 - *Crime reports at point level, demographics at block level.*
- The existing administrative boundaries are **human-defined** and **static**.
 - *Chicago administrative boundaries are defined by university scholars and keep being used for over 100 years.*
- **Task dependent partition:** for different prediction tasks, the partition may not be the same.
 - *In State College the Beaver stadium is a hotspot for traffic but not for crime incidents.*

Problem definition: task-specific region partition

Input:

- Target property observed at **tract** level: \mathbf{y} .
- Other properties (e.g., demographics) available for all **tracts**: \mathbf{x} .
- Spatial adjacency of all **tracts**: \mathbf{G} .

Output:

- Partition **tracts** into **community areas** with \mathbf{Y} and \mathbf{X} , such that

Given a fixed partition Z , the optimal task function f can be easily learned.

The challenge lies in searching through the partition space.

$$\arg \min_{Z, f} \sum_{j=1}^m \left(\underbrace{\|Y_j - f(X_j)\|_2}_{\text{Task prediction error}} + \underbrace{G(Z_j)}_{\text{Constraint on partition, e.g. variance of population}} \right) \quad (\text{Eq. 1})$$

Z : partition. f : task function on **community**
(mapping from tracts to community) (regression model, etc.)

Task prediction error

Constraint on partition, e.g.
variance of population

A partition over City T is denoted as $Z = \{Z_1, Z_2, \dots, Z_m\}$, satisfying the following:

- (1) (subset) $\forall j, Z_j \subset T$; (2) (non-overlapping) $\forall p, q, Z_p \cap Z_q = \emptyset$; (3) (completeness) $\bigcup \{Z_j\} = T$;
(4) (spatial-continuity) $\forall j, P_j$ defines a polygon with exact one connected component.

This problem is challenging

1. It's difficult to calculate the maximum likelihood of objective function.
 - a. Gradients on discrete variables are non available.
 - b. Region properties (features and target variable) and model coefficients are high dimensional
 - c. region properties change simultaneously when we change the region partition
 - d. Complicated constraints should be met.
2. It is a NP-hard combinatorial optimization problem.
 - a. Brute Force Solution: Enumerate all the possible combinations
 - $O(M^N)$ - M is the number of tracts, N is the number of communities

Solution: Techniques like Simulated Annealing to approximate global optimization

- Approximate the global optimum with the Markov Chain Monte Carlo (MCMC) method

Markov Chain Monte Carlo (Metropolis-Hastings)

Motivation: Obtaining a sequence of random samples from a distribution for which direct sampling is difficult. It constructs a Markov chain that will ultimately converge through stochastic sampling. We only **care about the last state**.

In our case:

- state space is all possible partitions Z
- quality measure $\mathcal{F}(Z)$, a partition Z with lower value is more likely to be optimal
- **proposal function $q(Z'|Z)$** , defines the transition probability
- $p(Z)$ is the Boltzmann distribution $p(Z) = \frac{e^{-\mathcal{F}(Z)/T}}{P}$
- Stopping criteria:
 - Number of iterations/the standard deviation of $\mathcal{F}(Z)$

Algorithm 1 MCMC method to search \mathcal{Z} .

```
1:  $\mathcal{Z} \leftarrow \mathcal{Z}_0$ 
2: while not satisfies stopping criteria do
3:   Sample  $u \leftarrow \mathcal{U}_{[0,1]}$ 
4:   Sample  $\mathcal{Z}' \leftarrow q(\mathcal{Z}'|\mathcal{Z})$ 
5:    $\gamma = \min \left[ 1, \frac{p(\mathcal{Z}')q(\mathcal{Z}|\mathcal{Z}')}{p(\mathcal{Z})q(\mathcal{Z}'|\mathcal{Z})} \right]$ 
6:   if  $u < \gamma$  then
7:      $\mathcal{Z} \leftarrow \mathcal{Z}'$ 
8:   end if
9: end while
```


Naive MCMC

- Start from administrative boundary
- **Randomly** select one tract to **flip its assignment**.
- Decide whether to accept the new partition with training error on prediction task f .

$$q(\mathcal{Z}'|\mathcal{Z}) = q(\mathcal{Z}|\mathcal{Z}') = \frac{1}{|\mathcal{T}_b| \cdot |\text{Adjacent}(t_i)|}$$

$$\gamma = \min\left[1, \frac{p(\mathcal{Z}')}{p(\mathcal{Z})}\right]$$

This could last a long time - since it's randomly select one tract to flip




Flip the assignment of a tract.

Guided MCMC - Softmax MCMC

- Start from administrative boundary
- Select tract from Community with the highest prediction error** (according to a softmax function), then **flip its assignment**.
- Decide whether to accept the new partition with training error on prediction task f .

$$Q(Z_j) = \frac{\exp(\|Y_j - f(X_j)\|_2)}{\sum_{k=1}^m \exp(\|Y_k - f(X_k)\|_2)}$$

$$q(\mathcal{Z}' | \mathcal{Z}) = \frac{Q(Z_j)}{|Z_j \cap \mathcal{T}_b| \cdot |\text{Adjacent}(t_i)|}$$

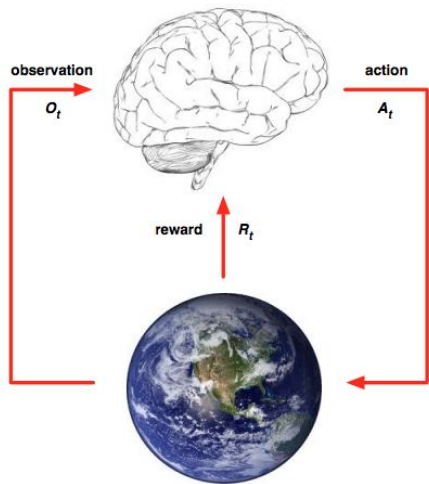


Flip the assignment of a tract $T2$.

Tract $T2$ is selected because the green community has the highest prediction error.



MCMC with reinforcement learning (RL)

- Start from administrative boundary
- **Instead of pre-define sampling strategy, learn where to sample with RL, then flip its assignment.**



At each step t the agent:

- Executes action A_t
- Receives observation O_t
- Receives scalar reward R_t



Tract T3 is selected because based on past experience, T3 is more likely to decrease \mathcal{F} most

MCMC with reinforcement learning (RL)

- Start from administrative boundary
- **Instead of pre-define sampling strategy, learn where to sample with RL, then flip its assignment.**

Challenges:

- Huge state space
- Large and dynamic action space
- Training overhead is high

Learning

$$Q(\mathcal{Z}, \langle t^k, Z^k \rangle) = \Delta \mathcal{F} + \delta \cdot \sum_{a \in \{\langle t^{k+1}, Z^{k+1} \rangle\}} Q(\mathcal{Z}', a)$$

\mathcal{Z} : current partition state

$\langle t, Z \rangle$: action (a tract to flip)

- Sampling: $\arg \max_{\langle t, Z \rangle} Q(\mathcal{Z}, \langle t, Z \rangle)$

Q is a neural network
 δ is set to 0

Take the max of 30
random samples

Results

Method	MAE	
	Crime	House price
Admin	1715.91	31.29
Agglomerative	72201.00	50.34
K-means	2887.83	32.40
Spectral	1440.57	29.66
Naive	1073.42(81.93)	25.73(2.76)
Softmax	1041.68(76.75)	27.13(2.98)
DQN	746.13(154.19)	25.16(1.30)


MAE of two prediction tasks.

Dataset:

[1] <https://data.cityofchicago.org/>

[2] <https://www.zillow.com/>

Clustering first, then do prediction




Convergence of three proposed methods.

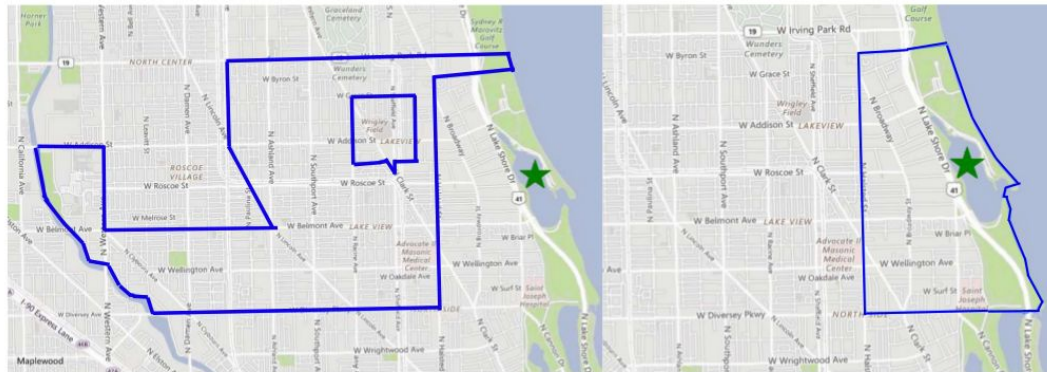
Case study: Community #6



Community #6 from
administrative boundary



DQN partition around
Community #6



Zillow's self-defined region (Zillow.com: real estate website)

The green star marks the location of **Belmont Harbor** (one of Chicago's largest boating areas).

DQN community #7: contain leisure destinations such as the Lincoln Park Zoo, and numerous beach areas.

Suggestions?
